

Predicting Backorders using Machine Learning

Lu Yang

a. Abstract

Backorder is a common supply chain problem which can cause many bad effects. Multiple suppliers and limited time that customer are willing to wait make this problem even more difficult to solve. In this paper, two machine learning models are used to build a predictive model for monitoring backorders. They are decision tree and random forest. Sampling and dummy variables are employed in these two models.

b. Introduction

1. Question: Can you predict product backorders?

In other words, can you predict backorder risk for products based on historical data?

2. Data Source:

- Can you predict product backorders?

<https://www.kaggle.com/tiredgeek/predict-bo-trial>

This dataset includes training data and test data.

Training dataset contains the descriptive data for the 8 weeks prior to the week we are trying to predict. These data was recorded weekly.

- Columns:

- 1) Predictor variables:

sku - Random ID for the product

national_inv - Current inventory level for the part

lead_time - Transit time for product (if available)

in_transit_qty - Amount of product in transit from source

forecast_3_month - Forecast sales for the next 3 months

forecast_6_month - Forecast sales for the next 6 months

forecast_9_month - Forecast sales for the next 9 months

sales_1_month - Sales quantity for the prior 1 month time period

sales_3_month - Sales quantity for the prior 3 month time period

sales_6_month - Sales quantity for the prior 6 month time period

sales_9_month - Sales quantity for the prior 9 month time period

min_bank - Minimum recommend amount to stock

potential_issue - Source issue for part identified

pieces_past_due - Parts overdue from source

perf_6_month_avg - Source performance for prior 6 month period

perf_12_month_avg - Source performance for prior 12 month period

local_bo_qty - Amount of stock orders overdue

deck_risk - Part risk flag

oe_constraint - Part risk flag

ppap_risk - Part risk flag

stop_auto_buy - Part risk flag

rev_stop - Part risk flag

2) Target variable:

went_on_backorder - Product actually went on backorder

3. Backgrounds:

1) What is a backorder?

Backorders are products that are temporarily out of stock, but will ship to customers when it is available.

In other word, a backorder generally indicates that customer demand for a product or service exceeds a company's capacity to supply it.

2) Backorder is a big problem.

- Huge number of sales orders and different suppliers where the out of stock items from increase complexity to the workload.
- Customers may not have the patience to wait for items. It may lead to lost sales and low customer satisfaction.

3) Way to handle it:

In order to monitor backorder, we can use machine learning to identify products at risk of backorders. Sales person can prepare for these products in advance.

Machine learning algorithm: Naive Bayes, Support vector machine, Decision tree, Random forest, K nearest neighbors

c. Code with Documentation

EDA:

1. Load and View my data

Load my data:

```
In [22]: df = pd.read_csv('C:\\Users\\LucyYang\\Jupyter\\Kaggle_Training_Dataset_v2.csv')
df.head()
```

E:_LastSemester\\2018Spring\\DataScience\\Tool\\Anaconda3\\lib\\site-packages\\IPython\\core\\interactiveshell.py:2698: DtypeWarning: Columns (0) have mixed types. Specify dtype option on import or set low_memory=False.

```
Out[22]:
```

	sku	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month	sales_1_month	sales_3_month	sales_6_month	...	
0	1026827	0.0	NaN	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
1	1043384	2.0	9.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
2	1043696	2.0	NaN	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
3	1043852	7.0	8.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	
4	1044048	8.0	NaN	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	

5 rows × 23 columns

View basic info:

```

In [23]: df.shape
Out[23]: (1687861, 23)

In [24]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1687861 entries, 0 to 1687860
Data columns (total 23 columns):
sku                1687861 non-null object
national_inv       1687860 non-null float64
lead_time          1596967 non-null float64
in_transit_qty     1687860 non-null float64
forecast_3_month   1687860 non-null float64
forecast_6_month   1687860 non-null float64
forecast_9_month   1687860 non-null float64
sales_1_month      1687860 non-null float64
sales_3_month      1687860 non-null float64
sales_6_month      1687860 non-null float64
sales_9_month      1687860 non-null float64
min_bank           1687860 non-null float64
potential_issue     1687860 non-null object
pieces_past_due     1687860 non-null float64
perf_6_month_avg    1687860 non-null float64
perf_12_month_avg   1687860 non-null float64
local_bo_qty        1687860 non-null float64
deck_risk           1687860 non-null object
oe_constraint       1687860 non-null object
ppap_risk           1687860 non-null object
stop_auto_buy       1687860 non-null object
rev_stop            1687860 non-null object
went_on_backorder   1687860 non-null object
dtypes: float64(15), object(8)
memory usage: 296.2+ MB

```

View categorical variables:

```

In [25]: df.deck_risk.value_counts()
df.oe_constraint.value_counts()
df.ppap_risk.value_counts()
df.stop_auto_buy.value_counts()
df.rev_stop.value_counts()
df.went_on_backorder.value_counts()
df.potential_issue.value_counts()

Out[25]: No      1686953
Yes        907
Name: potential_issue, dtype: int64

```

2. Data cleaning

Check missing value:

1) Deal with column 'lead_time'

a) deal with column 'lead_time'


the 'leadtime' is a skew right distribution, so we replace the missing value with the mean value (In practice, for skewed distributions the most commonly reported typical value is the mean)

```

In [27]: df.lead_time.isnull().sum()/df.shape[0]
Out[27]: 0.059775692296754473

In [28]: plt.show()
df.lead_time.hist(figsize = (15, 20))

```



```
In [29]: df.lead_time.fillna(df.lead_time.mean(), inplace = True)
```

```
In [30]: df.lead_time.isnull().sum()
```

```
Out[30]: 0
```

2) deal with column 'perf_6_month_avg' (for there missing value is -99)

b) deal with column 'perf_6_month_avg' (for there missing value is -99)


the 'perf_6_monthavg' is a skew left distribution, so we replace the missing value with the mean value (In practice, for skewed distributions the most commonly reported typical value is the mean)

```
In [31]: df.perf_6_month_avg.value_counts(dropna = False).values[2]/df.shape[0]
```

```
Out[31]: 0.076711338618131841
```

```
In [32]: df.perf_6_month_avg.replace(-99.00, np.nan, inplace = True)
```

```
In [33]: plt.show()
df.perf_6_month_avg.hist(figsize = (15, 20))
```



```
In [34]: df.perf_6_month_avg.fillna(df.perf_6_month_avg.mean(), inplace = True)
df.perf_6_month_avg.isnull().sum()

Out[34]: 0
```

3) deal with column 'perf_12_month_avg' (for there missing value is -99)

c) deal with column 'perf_12_month_avg' (for there missing value is -99)


the 'perf_12_monthavg' is a skew left distribution, so we replace the missing value with the mean value (In practice, for skewed distributions the most commonly reported typical value is the mean)

```
In [35]: df.perf_12_month_avg.value_counts(dropna = False).values[2]/df.shape[0]

Out[35]: 0.069710758001256035
```

```
In [36]: df.perf_12_month_avg.replace(-99.00, np.nan, inplace = True)
```

```
In [81]: plt.show()
df.perf_12_month_avg.hist(figsize = (15, 20))
```



Out[81]: <matplotlib.axes._subplots.AxesSubplot at 0x2394c484828>

```
In [38]: df.perf_12_month_avg.fillna(df.perf_12_month_avg.mean(), inplace = True)
df.perf_12_month_avg.isnull().sum()
```

Out[38]: 0

Check inappropriate value

check data type of each column

```
In [39]: df.dtypes
Out[39]: sku                object
national_inv            float64
lead_time              float64
in_transit_qty         float64
forecast_3_month       float64
forecast_6_month       float64
forecast_9_month       float64
sales_1_month          float64
sales_3_month          float64
sales_6_month          float64
sales_9_month          float64
min_bank              float64
potential_issue        object
pieces_past_due        float64
perf_6_month_avg       float64
perf_12_month_avg      float64
local_bo_qty          float64
deck_risk              object
oe_constraint          object
ppap_risk              object
stop_auto_buy          object
rev_stop               object
went_on_backorder      object
dtype: object
```

check each column whether they have inappropriate value


```
In [40]: df.national_inv.value_counts(dropna = False)
df.in_transit_qty.loc[df['in_transit_qty'] < 0]
df.forecast_3_month.loc[df['forecast_3_month'] < 0]
df.forecast_6_month.loc[df['forecast_6_month'] < 0]
df.forecast_9_month.loc[df['forecast_9_month'] < 0]
df.sales_1_month.loc[df['sales_1_month'] < 0]
df.sales_3_month.loc[df['sales_3_month'] < 0]
df.sales_6_month.loc[df['sales_6_month'] < 0]
df.sales_9_month.loc[df['sales_9_month'] < 0]
df.min_bank.loc[df['min_bank'] < 0]
df.pieces_past_due.loc[df['pieces_past_due'] < 0]
df.local_bo_qty.loc[df['local_bo_qty'] < 0]

Out[40]: Series([], Name: local_bo_qty, dtype: float64)
```


3. Data distribution

1) Each column distribution

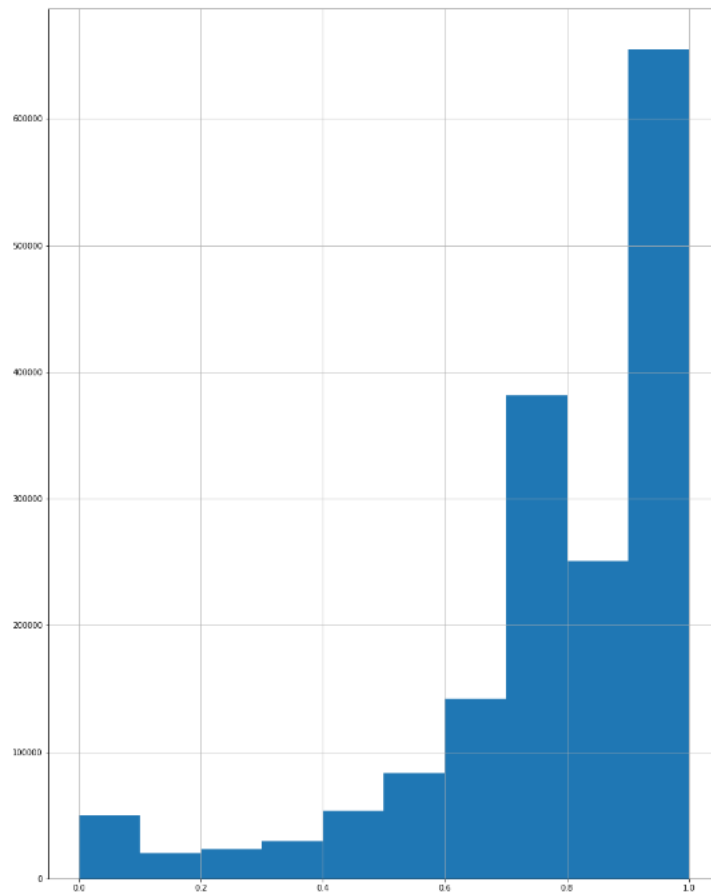
forecast_3_month is a skew right distribution. Most values are gathered in [0, 20000] (the mean is greater than the median.)



lead_time is a skew right distribution. Most values are gathered in [0, 60] (the mean is greater than the median.)



perf_6_month_avg is a skew left distribution. The range is [0, 1] (the mean is less than the median.)



2) Summary statistics

Question 3.2: What are the summary statistics?

Answer as follow:

```
In [50]: df.describe()
```

Out[50]:

	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month	sales_1_month	sales_3_month	sales_6_month	sale:
count	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.687860e+06	1.6
mean	4.961118e+02	7.872267e+00	4.405202e+01	1.781193e+02	3.449867e+02	5.063644e+02	5.592607e+01	1.750259e+02	3.417288e+02	5.2
std	2.961523e+04	6.841885e+00	1.342742e+03	5.026553e+03	9.795152e+03	1.437892e+04	1.928196e+03	5.192378e+03	9.613167e+03	1.4
min	-2.725600e+04	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.0
25%	4.000000e+00	4.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.0
50%	1.500000e+01	8.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00	2.000000e+00	4.0
75%	8.000000e+01	8.000000e+00	0.000000e+00	4.000000e+00	1.200000e+01	2.000000e+01	4.000000e+00	1.500000e+01	3.100000e+01	4.7
max	1.233440e+07	5.200000e+01	4.894080e+05	1.427612e+06	2.461360e+06	3.777304e+06	7.417740e+05	1.105478e+06	2.146625e+06	3.2


3) Outliers

Question 3.3: Are there anomalies/outliers?

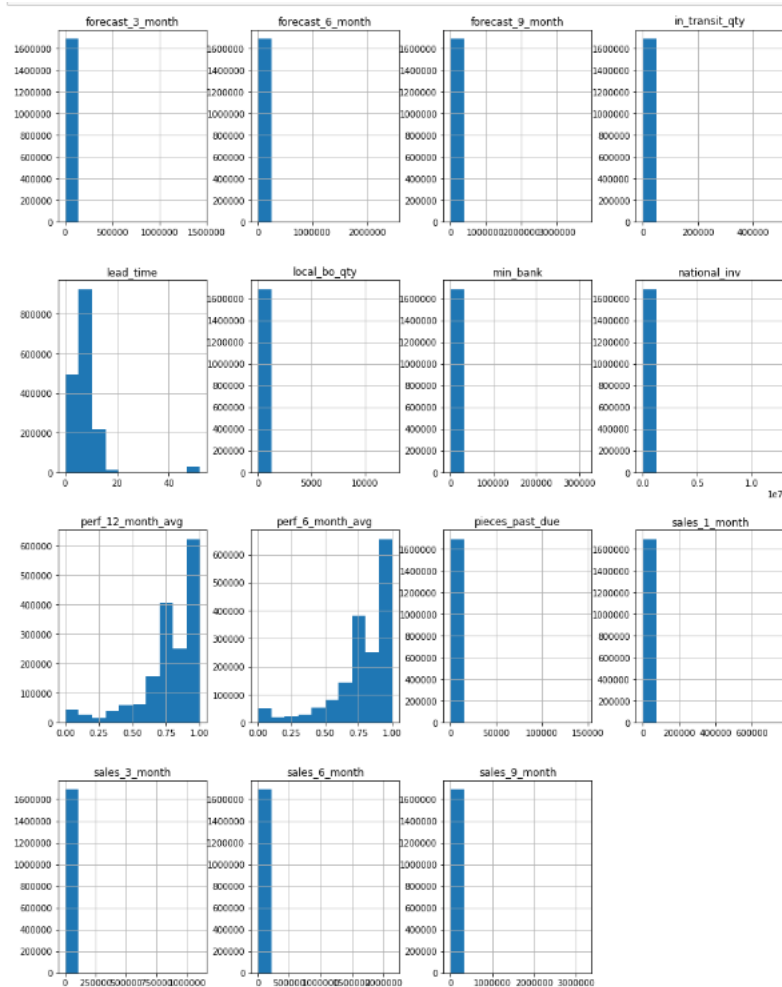
Answer as follow:

1. national_inv has most number of outliers and they are very discrete.
2. (forecast_3_month, forecast_6_month, forecast_9_month), (sales_3_month, sales_6_month, sales_9_month) are two group which outliers are distributed more and more discrete.
3. in_transit_qty, min_bank, pieces_past_due have a few number of outliers
4. leadtime, perf_6_month_avg, perf_12_month_avg, local_bo_qty are gathered.

```
In [52]: plt.show()  
df.plot(kind='box', figsize=(30, 50))
```



4) Plot each column



Summary

- Most columns have a very clustered and representative range.
- Lead time, perf_6_month_avg, perf_12_month_avg cannot find a typical value to represent these columns

5) Columns Correlating

Are any of the columns correlated?

These 3 group are correlated columns:


1. forecast_3_month, forecast_6_month, forecast_9_month
2. sales_3_month, sales_6_month, sales_9_month
3. perf_6_month_avg, perf_12_month_avg

```
In [55]: df.corr()
```

```
Out[55]:
```

	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month	sales_1_month	sales_3_month	sales_6_mon
national_inv	1.000000	0.003318	0.098238	0.078199	0.079744	0.078948	0.147449	0.192605	0.225067
lead_time	0.003318	1.000000	-0.006844	-0.008008	-0.008511	-0.008736	-0.005609	-0.006758	-0.007097
in_transit_qty	0.098238	-0.006844	1.000000	0.662648	0.687768	0.679152	0.619270	0.698417	0.689908
forecast_3_month	0.078199	-0.008008	0.662648	1.000000	0.990490	0.977337	0.684494	0.781178	0.835585
forecast_6_month	0.079744	-0.008511	0.687768	0.990490	1.000000	0.994945	0.701770	0.808755	0.868099
forecast_9_month	0.078948	-0.008736	0.679152	0.977337	0.994945	1.000000	0.716367	0.829911	0.891884
sales_1_month	0.147449	-0.005609	0.619270	0.684494	0.701770	0.716367	1.000000	0.918548	0.867479
sales_3_month	0.192605	-0.006758	0.698417	0.781178	0.808755	0.829911	0.918548	1.000000	0.975594
sales_6_month	0.225067	-0.007097	0.689908	0.835585	0.868099	0.891884	0.867479	0.975594	1.000000
sales_9_month	0.239613	-0.007239	0.659372	0.825539	0.858253	0.881894	0.815959	0.929491	0.971884
min_bank	0.399969	-0.007088	0.749974	0.725042	0.738553	0.735891	0.756137	0.856017	0.837147
pieces_past_due	0.030677	-0.001500	0.167460	0.361214	0.363147	0.366001	0.249526	0.304565	0.323585
perf_6_month_avg	0.002869	-0.009766	0.004570	0.006739	0.007503	0.007786	0.006323	0.006938	0.007111
perf_12_month_avg	0.002037	-0.007482	0.001611	0.003357	0.003717	0.003768	0.002058	0.002091	0.001949
local_bo_qty	0.014887	-0.001255	0.066612	0.039419	0.039724	0.039732	0.066188	0.071030	0.057711

```
In [131]: plt.show()
plt.matshow(df.corr())
```



```
Out[131]: <matplotlib.image.AxesImage at 0x2398469b780>
```

Analyze data:

1. Data transform

Converting categorical data into number:

```
In [57]: df_with_dummies = pd.get_dummies(df, columns = ['potential_issue', 'deck_risk', 'oe_constraint', 'ppap_risk', 'stop_auto'])
df_with_dummies.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1687860 entries, 0 to 1687859
Data columns (total 29 columns):
sku                    1687860 non-null object
national_inv          1687860 non-null float64
lead_time             1687860 non-null float64
in_transit_qty        1687860 non-null float64
forecast_3_month      1687860 non-null float64
forecast_6_month      1687860 non-null float64
forecast_9_month      1687860 non-null float64
sales_1_month         1687860 non-null float64
sales_3_month         1687860 non-null float64
sales_6_month         1687860 non-null float64
sales_9_month         1687860 non-null float64
min_bank             1687860 non-null float64
pieces_past_due       1687860 non-null float64
perf_6_month_avg      1687860 non-null float64
perf_12_month_avg     1687860 non-null float64
local_bo_qty         1687860 non-null float64
went_on_backorder     1687860 non-null object
potential_issue_No    1687860 non-null uint8
potential_issue_Yes   1687860 non-null uint8
deck_risk_No         1687860 non-null uint8
deck_risk_Yes        1687860 non-null uint8
oe_constraint_No      1687860 non-null uint8
oe_constraint_Yes     1687860 non-null uint8
ppap_risk_No          1687860 non-null uint8
ppap_risk_Yes        1687860 non-null uint8
stop_auto_buy_No     1687860 non-null uint8
stop_auto_buy_Yes    1687860 non-null uint8
rev_stop_No          1687860 non-null uint8
rev_stop_Yes         1687860 non-null uint8
dtypes: float64(15), object(2), uint8(12)
memory usage: 251.1+ MB
```

Transform data for model:

```
In [58]: df_with_dummies.shape
Out[58]: (1687860, 29)

In [59]: x_transform = df_with_dummies.drop(['sku', 'went_on_backorder'], axis=1)

In [60]: x_transform.head()
Out[60]:
```

	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month	sales_1_month	sales_3_month	sales_6_month	sales_9_month
0	0.0	7.872267	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	2.0	9.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	2.0	7.872267	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	7.0	8.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	8.0	7.872267	0.0	0.0	0.0	0.0	0.0	0.0	0.0	4.0

5 rows × 27 columns

```
In [61]: lb = pre.LabelBinarizer()
y_transform = lb.fit_transform(df_with_dummies.went_on_backorder)
```

Sampling into train data and test data

```
In [62]: # split dataset into inputs and outputs
values = df_with_dummies.values
x = x_transform
y = y_transform.ravel()
x = x.astype('int')
y = y.astype('int')
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=6)
```

2. Default decision tree

1) decision tree

```
In [63]: clf_tree = tree.DecisionTreeClassifier()
         clf_tree.fit(x_train, y_train)

Out[63]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

show result

```
In [64]: y_pred_tree = clf_tree.predict(x_test)
         tree_result = accuracy_score(y_test, y_pred_tree.round())

Out[64]: 0.99144775040583932
```

show confusion matrix

```
In [68]: confusion_matrix(y_test, y_pred_tree.round())

Out[68]: array([[333896,   1366],
                [ 1521,    789]], dtype=int64)
```

3. Decision tree(Max Depth = 3)

Build a new model to avoid overfitting

```
In [69]: clf_tree_new = tree.DecisionTreeClassifier(max_depth = 3)
         clf_tree_new.fit(x_train, y_train)
         y_pred_tree_new = clf_tree_new.predict(x_test)
         tree_result_new = accuracy_score(y_test, y_pred_tree_new.round())

Out[69]: 0.99315701539227186
```

show confusion matrix

```
In [71]: confusion_matrix(y_test, y_pred_tree_new.round())

Out[71]: array([[335262,    0],
                [ 2310,    0]], dtype=int64)
```

Draw the decision tree:

Draw the decision tree

```
In [38]: dot_data = tree.export_graphviz(clf_tree_new, out_file=None)
         graph = graphviz.Source(dot_data)
         graph.render("backorder")

In [ ]: dot_data = tree.export_graphviz(clf_tree_new, out_file=None,
                                         feature_names=list(x_transform),
                                         class_names=["IfBackorder"],
                                         filled=True, rounded=True,
                                         special_characters=True)
         graph = graphviz.Source(dot_data)
         graph
```

4. Random forest

2) Random Forest

```
In [72]: clf = RandomForestClassifier(n_jobs=2, random_state=0)
         clf.fit(x_train, y_train)

Out[72]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=2,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)

show result

In [74]: y_pred_clf = clf.predict(x_test)
         accuracy_score(y_test, y_pred_clf.round())

Out[74]: 0.99426196485490503

show confusion matrix

In [75]: confusion_matrix(y_test, y_pred_clf.round())

Out[75]: array([[335118, 144],
               [ 1793,  517]], dtype=int64)
```

d. Results

● Summary of EDA:

1. Overall:

1) This dataset is a very big dataset which have 1687861 rows and 23 columns. It means that I need build a large model for this dataset.

2) This dataset has many categorical variables which I need to convert it to numerical variables.

3) The target variable is categorical. So I should choose classification mode to analyse it.

2. Model recommendation:

1) decision tree

2) random forest

3. Data cleaning:

1) This dataset has some missing value: lead_time has NaN; perf_6_month_avg and perf_6_month_avg have -99 as missing value. I use mean to replace the missing value because they are skewed distribution which use mean as typical value.

2) This dataset doesn't have inappropriate values

4. Data distribution:

Most columns are skewed distribution but have a very clustered and representative range.

1) national_inv has the most number of outliers and are most discrete.

2) (forecast_3_month, forecast_6_month, forecast_9_month), (sales_3_month, sales_6_month, sales_9_month) are two groups which are more and more discrete.

3) in_transit_qty, min_bank, pieces_past_due have a few number of outliers and not very discrete.

4) leadtime, perf_6_month_avg, perf_12_month_avg, local_bo_qty are gathered.

5) Lead time, perf_6_month_avg, perf_12_month_avg cannot find a typical value to represent these columns

5. Column correlated:

(forecast_3_month, forecast_6_month, forecast_9_month), (sales_3_month, sales_6_month, sales_9_month), (perf_6_month_avg, perf_12_month_avg) are column correlated.

● Model results:

1. Accuracy_Score

Model	Accuracy_Score
Default Decision Tree	0.99144775040583932
Decision Tree (max_depth = 3)	0.99315701539227186
Random Forest	0.99426196485490503

2. Confusion Matrix

Default Decision Tree

333896	1366
1521	789

Decision Tree (max_depth = 3)

335262	0
2310	0

Random Forest

335118	144
1793	517

e. Discussion

- To build a predictive model for this classification problem, I try to use SVM model firstly. Support vector machine is a supervised learning model used for classification. However, it runs too slow to build a model with 1687861 records.
- Next model I used is decision tree. Decision tree is a decision support tool that uses a tree-like graph or model of decisions and their possible consequences, including chance event outcomes, resource costs, and utility. I used default parameters this time and got a good result. The accuracy score of this prediction is 0.991.
- To avoid overfitting, I set max_depth = 3 to build a new decision tree. This time I got a better accuracy score: 0.993.
- In order to get a better predictive model, I tried random forest at the end. Random forest is an ensemble learning method for classification, regression and other tasks, that operate by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes. I got the best result among these 3 modes: 0.994.
- **Conclusion**
 - 1) One of the effective way to avoid overfitting is set max depth for the decision tree model.
 - 2) Random forest is an optimized version of decision tree which can predict the result more accuracy than decision tree.
 - 3) SVM is not very suitable for a very large dataset.

f. References

1. Matt, Dancho (2017). Sales Analytics: How to Use Machine Learning to Predict and Optimize Product Backorders. Retrieved from http://www.business-science.io/business/2017/10/16/sales_backorder_prediction.html
2. Yuqi, Li (2017). AALBORG UNIVERSITY Backorder Prediction Using Machine Learning For Danish Craft Beer Breweries. Retrieved from

http://projekter.aau.dk/projekter/files/262657498/master_thesis.pdf

3. Rodrigo, Santis (2017). *Predicting Material Backorders in Inventory Management using Machine Learning*. Retrieved from

https://www.researchgate.net/publication/319553365_Predicting_Material_Backorders_in_Inventory_Management_using_Machine_Learning

4. Jason, Brownlee (2017). *How to Handle Missing Data with Python*. Retrieved from

<https://machinelearningmastery.com/handle-missing-data-python/>

5. FastML (2017). *Converting categorical data into numbers with Pandas and Scikit-learn*. Retrieved from

<http://fastml.com/converting-categorical-data-into-numbers-with-pandas-and-scikit-learn/>

6. Mode (2017). *Python Histograms, Box Plots, & Distributions*. Retrieved from

<https://community.modeanalytics.com/python/tutorial/python-histograms-boxplots-and-distributions/>